

## Bibliography

1. Bowman, S. R., Angeli, G., Potts, C., & Manning, C. D. (2015).  
A Large Annotated Corpus for Learning Natural Language Inference.  
EMNLP 2015, pp. 632–642.
2. Williams, A., Nangia, N., & Bowman, S. R. (2018).  
A Broad-Coverage Challenge Corpus for Sentence Understanding through Inference.  
NAACL-HLT 2018, pp. 1112–1122.
3. Nie, Y., Williams, A., Dinan, E., Bansal, M., Weston, J., & Kiela, D. (2020).  
Adversarial NLI: A New Benchmark for Natural Language Understanding.  
ACL 2020, pp. 4885–4901.
4. Arakelyan, E., Liu, Z., & Augenstein, I. (2024).  
Semantic Sensitivities and Inconsistent Predictions: Measuring the Fragility of NLI Models.  
EACL 2024, Volume 1: Long Papers, pp. 432–444.
5. McCoy, R. T., Pavlick, E., & Linzen, T. (2019).  
Right for the Wrong Reasons: Diagnosing Syntactic Heuristics in Natural Language Inference.  
ACL 2019, pp. 3428–3448.
6. Glockner, M., Shwartz, V., & Goldberg, Y. (2018).  
Breaking NLI Systems with Sentences that Require Simple Lexical Inferences.  
ACL 2018, Short Papers, pp. 650–655.
7. Naik, A., Ravichander, A., Sadeh, N., Rose, C., & Neubig, G. (2018).  
Stress Test Evaluation for Natural Language Inference.  
COLING 2018, pp. 2340–2353.
8. Richardson, K., Hu, H., Moss, L. S., & Sabharwal, A. (2020).  
Probing Natural Language Inference Models through Semantic Fragments.  
AAAI 2020.
9. Srikanth, N., Carpuat, M., & Rudinger, R. (2024).  
How Often Are Errors in Natural Language Reasoning Due to Paraphrastic Variability?

Transactions of the Association for Computational Linguistics, 12, pp. 1143–1162.

10. Zgreabă, M., Deoskar, T., & Abzianidze, L. (2025).

MERGE: Minimal Expression-Replacement Generalization Test for Natural Language Inference.

arXiv preprint, arXiv:2510.24295.

11. Wu, W., & Last, M. (2025).

Transitive Self-Consistency Evaluation of NLI Models without Gold Labels.

EMNLP 2025, pp. 22626–22642.

12. Papernot, N., & McDaniel, P. (2018).

Deep k-Nearest Neighbors: Towards Confident, Interpretable and Robust Deep Learning.

arXiv preprint, arXiv:1803.04765.

13. Ulanovski, T., Blyachman, E., & Bechler-Speicher, M. (2026).

Improving LLM Final Representations with Inter-Layer Geometry.

arXiv preprint, arXiv:2603.22665.

14. Arakelyan, E. (2025).

Reasoning Inconsistencies and How to Mitigate Them in Deep Learning.

PhD Thesis, Department of Computer Science, Faculty of Science, University of Copenhagen.